



OpenAI Forward Deployed Engineer (FDE)

Interview process - candidate preparation guide

UNOFFICIAL CANDIDATE PREP | AUG 2026

OpenAI FDEs work close to strategic customers: clarifying ambiguous problems, designing systems, shipping production code, driving rollout and adoption, and feeding field learning back into Product and Research.

OpenAI publishes a general interview framework, not one fixed FDE loop. Official assessments may include pair coding, take-home projects and technical tests; final interviews are typically a 4-6 hour block with 4-6 people over 1-2 days. The stages below combine official guidance with recent candidate-reported FDE patterns.

+ Typical FDE Interview Loop

Order, duration and tool policy can vary by team, level, geography and candidate. Confirm the exact loop with your recruiter.

Introductory / Recruiter Screen

FOCUS

Background, motivation, mission fit and whether your experience matches embedded customer-facing engineering.

Reported: ~30 mins

WHAT TO EXPECT

- Discuss work/academic experience, motivations, goals, and why OpenAI / why FDE.
- Explain one complex project clearly enough for a non-specialist to follow.

Skills Assessment / Customer Take-home

FOCUS

Can you turn a vague customer need into a working, defensible system - not merely a demo?

Team-dependent
Reported build: ~1 week

WHAT TO EXPECT

- Officially, skills assessments may include pair coding, take-home projects, technical tests, or more than one assessment.
- Recent FDE reports describe working code + a running app + a walkthrough, followed by a live technical/customer deep-dive.
- Expect follow-ups on customer framing, build quality, adaptability, and rejected alternatives.

Coding Interview

FOCUS

Production-style implementation under time pressure: correctness first, then edge cases, readability and performance.

Reported: ~60 mins

WHAT TO EXPECT

- Build a working solution, then extend it as requirements change.
- Narrate trade-offs; show strong language fundamentals and clean abstractions.

System Design / LLM Deployment

FOCUS

Design model-backed systems that survive scale, ambiguity, failures, security boundaries, latency targets and cost constraints.

Reported: ~45-60 mins

WHAT TO EXPECT

- Discuss retrieval, evaluation, retries/idempotency, observability, throughput and failure recovery.
- Turn a vague customer goal into concrete architecture and measurable success criteria.

Final Interviews

FOCUS

OpenAI says final interviews focus on your area of expertise and are designed to stretch you beyond your comfort zone.

Official: 4-6 hours
4-6 people / 1-2 days

OFFICIAL + REPORTED

- Project / technical deep dive: present one complex system you built; defend ownership, architecture, scale, evaluation and trade-offs.
- Behavioral + customer judgment: AI perspective, customer empathy, conflict, cross-functional work, ownership and communicating technical limits clearly.

Decision / References

FOCUS

A recruiter may ask for references at this stage; exact reference practice varies by role and geography.

Official: generally within ~1 week

WHAT TO EXPECT

- OpenAI says candidates should generally hear within one week of final interviews.